

Churn Rates Minimization

Kelso Quan

May 15, 2019

Executive Summary

Churn rates is an important aspect to any financial tech or commercial company. Churning is when the customer cancels their subscription or ceases doing business with the company. The goal of this analysis is to minimize churn rates. In this analysis, complete case analysis was used. The number of observations was 27000 dropped to 8453 and the number of features used to find strong indicators of churn rate went from 31 to 28. The analysis used decision tree, random forest, discrete, real, and gentle Ada Boosting to model the best way to predict churn rates. Every model had very similar misclassification errors with the best model being a classification tree. The best features to indicate churn rates pertained to credit card information, rewards from credit cards, and user's financial information.

1 Introduction

Churn rates are of interest to the Financial Tech companies and commercial businesses that rely on subscriptions and sale of services. Churning means cancellation of a subscription which is bad for the company. All companies would like to minimize their customers' churn rates to maximize profit. Churn is a binary response with a customer either churned or not churned. The data set has 31 features and 27000 observations. Given the data, the analysis tries to minimize churn rates by determining which features contribute heavily to customers cancelling their subscriptions and which model is the best.

2 Methods

In this analysis, decision tree, discrete, real, and gentle Ada Boosting along with random forest were used. R/Rstudio was used to perform the analysis. For the boosting techniques, the data was divided into 2/3 training set and 1/3 test set with cross validation being used. For the random forest, the data was split 60% training set and 40% test set. tuneRF was a function used to determine how many features were optimal for the random forest model which turned out to be 5 with the minimal OOB error. For decision tree, cross validation was used to grow a large tree, then 1SE rule was applied to prune the tree into the best subtree. The data set had 31180 Na's. Due to heavy computational times, the analysis was done as a complete case analysis. After omitting observations with Na's, there were 8453 observations left in the data set. The analysis eliminated user (id), android_user, and app_web_user. User id was eliminated because it did not provide information in determining someone's churn rate. The variables android_user and app_web_user were also deleted because they were negatively correlated and binary. Being binary variables of negatively correlated variables signified that if the user was not an android_user meant the user was ios_user and if the user was not using application, then they were a web_user.

3 Results

Table 1: Table of the response variable

	count
!churn	5239
churn	3214

3.1 Exploratory Data Analysis

Table 1 shows the response was binary. About 60% of the users did not churn and 40% of the users did churn. This proportion was used to split the training and test data for the random forest model. In addition, there were several features that were so skewed in distributions that it was better for those variables to become binary features. These features include registered_phones, deposits, withdrawal, purchase partners, purchases, cc taken, cc_disliked, cc_liked, and cc_application begin. cc is short for credit card. Furthermore, the analysis had to define type 1 and type 2 errors. Type 1, false positive, is predicting that the customers had churned when actually they had not churned. Type 2, false negative, is predicting that the customers had not churned when really they had churned.

In table 2, there are strong correlations between cc_recommended, rewards_earned, and reward_rate. There was little to no correlation between age and credit_score. Age and credit_score were weakly, positively correlated to cc_recommended, rewards_earned, and reward_rate features.

Table 2: Correlation plot of non-binary, numerical features

	age	credit_score	cc_recommended	rewards_earned	reward_rate
age	1.00	0.02	0.14	0.13	0.12
credit_score	0.02	1.00	0.13	0.15	0.17
cc_recommended	0.14	0.13	1.00	0.89	0.85
rewards_earned	0.13	0.15	0.89	1.00	0.96
reward_rate	0.12	0.17	0.85	0.96	1.00

3.2 Model Fitting and Results

This analysis did a preliminary logistic regression to examine if it was a good model for the data. It did not perform well. Consequentially, the analysis progressed into a classification tree. The best split for the decision tree was purchase partners. That was followed by the features: rewards_earned, reward_rate, cc_recommended, and credit_score. But a single decision tree is not stable, so an ensemble of trees were created.

Table 3: Misclassifications Error of Methods Used

Method	False Positive	False Negative	Total Misclassifications
Single Tree	15.21%	46.03%	27.10%
Gentle AdaBoost	8.13%	68.3%	31%
Real AdaBoost	11.79%	58%	29.34%
Discrete AdaBoost	13.27%	55.6%	29.34%
Random Forest	18.49%	46.40%	29.10%


From table 3, the models have very similar misclassification errors. When looking at the total misclassification column, those error rates are strictly below 30%. The best model being the single classification tree with the lowest total misclassifications. Gentle AdaBoost had the lowest false positive rates among all the models.

The figures (a), (b), (c), (d) shows the variable importance when it comes to predicting churn rates. Financial factors such as cc_recommended, rejected_loans, and other credit card features are significant in predicting churn rates. Any features that were closer to the right side of the importance plot were considered influential to the users' churn rate. The variables concentrated on the left side of the variable importance plot were considered less significant when it came to determining whether a customer was going to churn or not.

4 Conclusion

In determining an user's churn rate, it is best to consider their credit card history and financial information. cc_recommended, purchase_partners, reward_rate, and rejected_loan were of the highest variable importance when determining whether a subscriber stays in business or cancels their subscription. The analysis was slightly biased because the training-test data set ratio was different between random forest and the three methods of AdaBoost. In addition, all the models had very similar results with the decision tree being the most interpretable. The other models required variable importance plots to indicate which features contributed most to churn rates. In future analysis, the training-test data set ratio will be the same. Another constraint was the sheer amount of Na's in the data set. A considerable portion of the data was lost when complete case analysis was used rather than imputation or another method to handle missing values. Financial Tech companies should consider an user's financial history and credit card information to determine whether that user is at risk of churning.


Quan/Documents/SchoolWork/Stat702/final/realVarPlot.png



(a) Random Forest

(b) Real AdaBoost

Quan/Documents/SchoolWork/Stat702/final/disVarPlot.png



(c) Gentle AdaBoost

(d) Discrete AdaBoost

5 References

Khan, Adnan
<https://www.kaggle.com/nanda1331/churn-rate-minimization>
stackoverflow.com
overleaf.com

Appendix: R Code

```
1 setwd("~/SchoolWork/Stat702/final")
2 churn.data = read.csv("churn.csv", header = T)
3 names(churn.data)
4 str(churn.data)
5 dim(churn.data)
6 sum(is.na(churn.data))
7 set.seed(702)
8
9
10 #Take log transform of age
11 churn.data$age<-log(churn.data$age)
12
13 #Set age as numerical
14 churn.data$age<-as.numeric(churn.data$age)
15
16 #Set response variable as a factor and add labels
17 churn.data$churn <- factor(churn.data$churn, levels=0:1, labels=c("!churn", "churn"))
18
19 #Set housing as factor
20 churn.data$housing <- factor(churn.data$housing)
21
22 #app_web_user correlated to web_user
23 #remove app_web_user for now we will modify based on step wise model building
24 cor.test(~app_web_user+web_user, data=churn.data)
25 churn.data$app_web_user<-NULL
26
27 #ios_user correlated to android_user
28 #remove android user for now, we will modify based on step wise model building
29 cor.test(~ios_user+android_user, data=churn.data)
30 churn.data$android_user<-NULL
31
32 #remove user feature
33 churn.data$user<-NULL
34
35 #Set zodiac_sign as factor
36 churn.data$zodiac_sign<-as.factor(churn.data$zodiac_sign)
37
38
39 #Change features to factor (5:9) and (11:13)
40
41 #Features 5:9 (deposits, withdrawal, purchase partners, purchases, cc taken)
42 churn.data[,5:9]=sapply(churn.data[,5:9], function(x) replace(x,x>0,1))
43
44 #Features 11:13 (cc_disliked,cc_liked,cc_application begin)
45 churn.data[,11:13]=sapply(churn.data[,11:13], function(x) replace(x,x>0,1))
46
47 #set deposits as factor
48 churn.data$deposits<-as.factor(churn.data$deposits)
49
50 #set withdrawal as factor
51 churn.data$withdrawal<-as.factor(churn.data$withdrawal)
52
53 #set purchases as factor
54 churn.data$purchases<-as.factor(churn.data$purchases)
55
56 #set purchases_partners as factor
57 churn.data$purchases_partners<-as.factor(churn.data$purchases_partners)
58
59 #set cc taken as factor
60 churn.data$cc_taken<-as.factor(churn.data$cc_taken)
61
62 #set cc_disliked as factor
63 churn.data$cc_disliked<-as.factor(churn.data$cc_disliked)
```

```

64 #set cc_liked as factor
65 churn.data$cc_liked<-as.factor(churn.data$cc_liked)
66
67 #set cc_application_begin as factor
68 churn.data$cc_application_begin<-as.factor(churn.data$cc_application_begin)
69
70 #set app_downloaded as factor
71 churn.data$app_downloaded<-as.factor(churn.data$app_downloaded)
72
73 #set web_user as factor
74 churn.data$web_user<-as.factor(churn.data$web_user)
75
76 #Set ios_user as factor
77 churn.data$ios_user<-as.factor(churn.data$ios_user)
78
79 #registered phone as binary and to factor
80 churn.data[,17]=sapply(churn.data[,17],function(x) replace(x,x>0,1))
81 churn.data$registered_phones<-as.factor(churn.data$registered_phones)
82
83 #waiting for loan as factor
84 churn.data$waiting_4_loan<-as.factor(churn.data$waiting_4_loan)
85
86 #cancelled loan as factor
87 churn.data$cancelled_loan<-as.factor(churn.data$cancelled_loan)
88
89 #Received loan as factor
90 churn.data$received_loan<-as.factor(churn.data$received_loan)
91
92 #set rejected loan as factor
93 churn.data$rejected_loan<-as.factor(churn.data$rejected_loan)
94
95 #set left for one month as factor
96 churn.data$left_for_one_month<-as.factor(churn.data$left_for_one_month)
97
98 #set left for more than twomonths to factor
99 churn.data$left_for_two_month_plus<-as.factor(churn.data$left_for_two_month_plus)
100
101 #set is_referred as factor
102 churn.data$is_referred<-as.factor(churn.data$is_referred)
103
104 churn.data <- na.omit(churn.data)
105 dim(churn.data)
106 summary(churn.data)
107 sum(is.na(churn.data))
108
109 set1<-churn.data[churn.data$churn=="churn",]
110 set0<-churn.data[churn.data$churn=="!churn",]
111
112 dim(set1) # 3214 28
113 dim(set0) # 5239 28
114 3214*2/3 # 2142.667
115 5239*2/3 #3492
116
117 training1 <- sample(1:3214,2143)
118 test1 <- (1:3214)[-training1]
119 sum((1:3214) == sort(c(training1,test1)))
120
121 training0 <- sample(1:5239,3492)
122 test0 <- (1:5239)[-training0]
123 sum((1:5239 == sort(c(training0,test0)))) #2788
124
125 train <- rbind(set1[training1,], set0[training0,])
126 test <- rbind(set1[test1,], set0[test0,])
127
128 numeric.var <- sapply(churn.data, is.numeric)

```

```

129 corr.matrix <- cor(churn.data[,numeric.var])
130 library(corrplot)
131 corrplot(corr.matrix, main = '\n\nCorrelation Plot for Numerical Variables')
132 dim(train)
133 dim(test)
134 set.seed(100)
135 logmod <- glm(churn ~ ., family = binomial(link = "logit"), data = churn.data)
136 summary(logmod)
137
138 logmodel <- glm(churn ~ ., family=binomial(link="logit"), data=train)
139 summary(logmodel)
140 logmodell <- glm(churn ~ housing+credit_score+purchases_partners+cc_taken+
141   cc_recommended+cc_liked+cc_application_begin+web_user+
142   ios_user+registered_phones+payment_type+received_loan+
143   rejected_loan+zodiac_sign+left_for_two_month_plus+rewards_earned+
144   reward_rate+is_referred,
145   family= binomial(link = "logit"), data = train)
146 summary(logmodell)
147 logmodel2 <- glm(churn ~ housing+credit_score+purchases_partners+cc_taken+
148   cc_recommended+cc_liked+cc_application_begin+web_user+
149   registered_phones+received_loan+
150   rejected_loan+left_for_two_month_plus+rewards_earned+
151   reward_rate+is_referred,
152   family= binomial(link = "logit"), data = train)
153 summary(logmodel2)
154
155 testmodel <- glm(churn ~ housing+credit_score+purchases_partners+cc_taken+
156   cc_recommended+cc_liked+cc_application_begin+web_user+
157   registered_phones+received_loan+
158   rejected_loan+left_for_two_month_plus+rewards_earned+
159   reward_rate+is_referred,
160   family= binomial(link = "logit"), data = test)
161 summary(testmodel)
162
163 #####
164 #Boosting
165 #####
166
167 #cross validation errors 5.3.3 islr
168 lmat <- matrix(c(0,1,4,0), nrow=2, byrow=T)
169
170 library(ada)
171
172 #Discrete adaboost using 10 fold xval, cp=0
173 default=rpart.control(xval=10,cp=0)
174 fitdis<-ada(churn ~ ., data=train, iter=50, loss="e", type="discrete", control=default)
175 print(fitdis) #training error 0.087
176
177
178 #misclassification rate = (false positive + false negative)/total
179 pred_val.2 <- predict(fitdis, newdata= test)
180 table.2<-table(test$churn, pred_val.2)
181 miss.rate2<-(table(test$churn, pred_val.2)[1,2]+table(test$churn, pred_val.2)[2,1])/length(
182   test$churn)
183 miss.rate2 #28.63% misclassification overall, discrete adaboost might overfit
184
185 #false positive and false negative rates
186 fp.rate2<-(table(test$churn, pred_val.2)[1,2])/(table(test$churn, pred_val.2)[1,2]+table(test
187   $churn, pred_val.2)[1,1])
188 fp.rate2 #16.77% false positive
189
190 fn.rate2<-(table(test$churn, pred_val.2)[2,1])/((table(test$churn, pred_val.2)[2,1])+table(
191   test$churn, pred_val.2)[2,2])
192 fn.rate2#47.03% false negatives

```

```

191
192 #variable importance plot
193 varplot(fitdis)
194 vip <- varplot(fitdis, plot.it=FALSE, type="scores")
195 round(vip, 4)
196
197 #Real adaboost
198 fitreal<-ada(churn~., data=train, iter=50, type="real",
199             control=rpart.control(maxdepth=2, cp=-1, minsplit=0))
200 fitreal
201 varplot(fitreal)
202 #misclassification rate = (false positive + false negative)/total
203 pred_val.3 <- predict(fitreal, newdata= test)
204 table.3<-table(test$churn, pred_val.3)
205 miss.rate3<-(table(test$churn, pred_val.3)[1,2]+table(test$churn, pred_val.3)[2,1])/length(
206   test$churn)
207 miss.rate3 #for real adaboost still at 29.84%
208
209 #false positive and false negative rates
210 fp.rate3<-(table(test$churn, pred_val.3)[1,2])/(table(test$churn, pred_val.3)[1,2]+table(test
211   $churn, pred_val.3)[1,1])
212 fp.rate3 #15.08% false positive
213
214 fn.rate3<-(table(test$churn, pred_val.3)[2,1])/((table(test$churn, pred_val.3)[2,1]+table(
215   test$churn, pred_val.3)[2,2])
216 fn.rate3#59.43% false negatives
217
218 #gentle adaboost
219 fitgen<-ada(churn~., data=train, test.x=test[, -1], test.y=test[, 1], iter=50,
220             type="gentle",
221             control=rpart.control(cp=-1, maxdepth=8))
222 (fitgen)
223 varplot(fitgen)
224 #misclassification rate = (false positive + false negative)/total
225 pred_val.4 <- predict(fitgen, newdata= test)
226 table.4<-table(test$churn, pred_val.4)
227 miss.rate4<-(table(test$churn, pred_val.4)[1,2]+table(test$churn, pred_val.4)[2,1])/length(
228   test$churn)
229 miss.rate4 #still at 29.77%
230
231 #false positive and false negative rates
232 fp.rate4<-(table(test$churn, pred_val.4)[1,2])/(table(test$churn, pred_val.4)[1,2]+table(test
233   $churn, pred_val.4)[1,1])
234 fp.rate4 #18.17% false positive
235
236 fn.rate4<-(table(test$churn, pred_val.4)[2,1])/((table(test$churn, pred_val.4)[2,1]+table(
237   test$churn, pred_val.4)[2,2])
238 fn.rate4#46.20% false negatives
239
240 # training the model
241 model_ada<-train(churn~., data=train, method='ada', tuneGrid=grid)
242 plot(model_ada)
243
244 set.seed(702)
245
246 ##random forest##
247 library(randomForest)
248 dim(train)
249 c.tune <- tuneRF(train[2:28], train$churn, ntreeTry=50, stepFactor=2,
250                 improve=0.05, trace=TRUE, plot=TRUE, dobest=FALSE, main = "mtry vs OOB

```



```

    error")
250 # mtry OOBError
251 # 3.OOB      3 0.2917480
252 # 5.OOB      5 0.2857143
253 # 10.OOB     10 0.2897959
254 # 5 is the best mtry
255 c.tune
256 plot(c.tune)
257 train.rf <- randomForest(churn~., data = train, mtry = 5, ntree = 501, norm.votes = F)
258 test.rf <- randomForest(churn~., data = test, mtry = 5, ntree = 501, norm.votes = F)
259 print(train.rf)
260 print(test.rf)
261 #test.rf OOB estimate of error rate: 29.42% (misclass rate)
262 # OOB estimate of error rate: 29.42%
263 # Confusion matrix:
264 # !churn churn class.error
265 # !churn 1415 332 0.1900401
266 # churn 497 574 0.4640523
267
268 #varImpPlot(train.rf, main = "Variable Importance of Churn Rates")
269 varImpPlot(test.rf, main = "Variable Importance of Churn Rates")
270 plot(train.rf, main = "MSE vs # of bootstrap Samples")
271 plot(test.rf, main="MSE vs. # of bootstrap Samples")
272 legend("center", legend = test.rf)
273 legend("topright",
274       legend=as.character(levels(mtcars$cyl)),
275       fill = rainbow(nlevels(mtcars$cyl)),
276       title = "cyl")
277
278 ### prediction ###
279 ind<-sample(2, nrow(churn.data), replace=T, prob=c(.6, .4)) #1/2 (.6/.4): training/testing
280 table(ind)
281 churn.rf<-randomForest(churn~., data=churn.data[ind==1,])
282 print(churn.rf) #OOB estimate of error rate: 27.7%
283 varImpPlot(churn.rf)
284 churn.pred <- randomForest(churn~., data = churn.data[ind == 2,])
285 print(churn.pred)
286 varImpPlot(churn.pred, main = "Variable Importance Plot")
287 # misclassification error rate: 28.67%
288 # Confusion matrix:
289 # !churn churn class.error
290 # !churn 1712 372 0.1785029
291 # churn 601 709 0.4587786
292
293 ##rpart comparisons###
294 library(rpart)
295 library(caret)
296 library("e1071")
297
298 my.control <- rpart.control(xval=10, cp=0)
299 tree <- rpart(churn~.,
300             data=churn.data, method="class", control=my.control)
301 plot(tree, margin = .1, uniform = T)
302 text(tree, use.n = T)
303 printcp(tree)
304 plotcp(tree)
305 0.76602 +0.012997 #0.779017 tree 16, 30 splits, 31 terminal nodes
306 besttree<-prune(tree, cp=0.0018)
307 plot(besttree, margin=.1)
308 text(besttree, use.n=T)
309 printcp(besttree)
310 besttree8 <- prune(besttree, cp = .004)
311 plot(besttree8, uniform = T)
312 text(besttree8, use.n = T)
313 printcp(besttree8)

```

```

314 pred.tree<- predict(besttree,newdata=churn.data[ind==2,], type='class')
315 table(observed=churn.data[ind==2,"churn"],predicted=pred.tree)
316 (588+350)/(1760 + 313 +562 + 703) #0.2810066
317
318 nsplits= besttree$cptable[,2]
319 test_error= besttree$cptable[,3]
320 xerror= besttree$cptable[,4]
321 xstd= besttree$cptable[,5]
322 plot(nsplits, test_error, type = 'l', xlab = "Number of splits", ylab = "Test error",
323      main = "Test error vs Number of splits")
324 plot(nsplits, test_error)
325 lines(nsplits, xerror, lty=3, col = "red")
326 lines(nsplits, xerror+xstd, lty=4, col = "blue")
327 lines(nsplits, xerror-xstd, lty=4, col = "blue")
328 legend(17,1, c("test error", "xerror", "+/- 1 xstd"), lty = c(1,3,4))
329 xtable(varImp(besttree))

```

Listing 1: Churn Rates